

# 0075\_k\_fold

April 20, 2018

## 1 Choosing between models with stratified k-fold validation

In previous example we have used multiple random sampling in order to obtain a better measurement of accuracy for models (repeating the model with different random training/test splits).

A more robust method is to use 'stratified k-fold validation'. In this method the model is repeated k times, so that all the data is used once, but only once, as part of the test set. This, alone, is k-fold validation. Stratified k-fold validation adds an extra level of robustness by ensuring that in each of the k training/test splits, the balance of outcomes represents the balance of outcomes in the overall data set. Most commonly 10 different splits of the data are used.

In this example we shall load up some data on treatment of acute stroke (data will be loaded from the internet). The model will try to predict whether patients are treated with a clot-busting drug. We will compare a number of different models using stratified k-fold validation.

```
In [35]: """Techniques applied:  
         1. Random Forests  
         2. Support Vector Machine (linear and rbf kernel)  
         3. Logistic Regression  
         4. Neural Network  
         """  
  
         # %% Load modules  
  
         import os  
         import numpy as np  
         import pandas as pd  
         from sklearn.model_selection import train_test_split  
         from sklearn.preprocessing import StandardScaler  
         from sklearn.ensemble import RandomForestClassifier  
         from sklearn.svm import SVC  
         from sklearn.linear_model import LogisticRegression  
         from sklearn.neural_network import MLPClassifier  
         from sklearn import datasets  
         from sklearn.model_selection import StratifiedKFold  
  
         # %% Function to calculate sensitivity and specificity  
         def calculate_diagnostic_performance(actual_predicted):
```

```

""" Calculate sensitivity and specificity.
Takes a Numpy array of 1 and zero, two columns: actual and predicted
Returns a tuple of results:
1) accuracy: proportion of test results that are correct
2) sensitivity: proportion of true +ve identified
3) specificity: proportion of true -ve identified
4) positive likelihood: increased probability of true +ve if test +ve
5) negative likelihood: reduced probability of true +ve if test -ve
6) false positive rate: proportion of false +ves in true -ve patients
7) false negative rate: proportion of false -ves in true +ve patients
8) positive predictive value: chance of true +ve if test +ve
9) negative predictive value: chance of true -ve if test -ve
10) Count of test positives

*false positive rate is the percentage of healthy individuals who
incorrectly receive a positive test result
* false negative rate is the percentage of diseased individuals who
incorrectly receive a negative test result

"""

actual_predicted = test_results.values
actual_positives = actual_predicted[:, 0] == 1
actual_negatives = actual_predicted[:, 0] == 0
test_positives = actual_predicted[:, 1] == 1
test_negatives = actual_predicted[:, 1] == 0
test_correct = actual_predicted[:, 0] == actual_predicted[:, 1]
accuracy = np.average(test_correct)
true_positives = actual_positives & test_positives
true_negatives = actual_negatives & test_negatives
sensitivity = np.sum(true_positives) / np.sum(actual_positives)
specificity = np.sum(true_negatives) / np.sum(actual_negatives)
positive_likelihood = sensitivity / (1 - specificity)
negative_likelihood = (1 - sensitivity) / specificity
false_postive_rate = 1 - specificity
false_negative_rate = 1 - sensitivity
positive_predictive_value = np.sum(true_positives) / np.sum(test_positives)
negative_predictive_value = np.sum(true_negatives) / np.sum(test_negatives)
positive_rate = np.mean(actual_predicted[:,1])
return (accuracy, sensitivity, specificity, positive_likelihood,
        negative_likelihood, false_postive_rate, false_negative_rate,
        positive_predictive_value, negative_predictive_value,
        positive_rate)

# %% Print diagnostics results
def print_diagnostic_results(results):
    # format all results to three decimal places
    three_decimals = ["%.3f" % v for v in results]

```

```

print()
print('Diagnostic results')
print('  accuracy:\t\t\t', three_decimals[0])
print('  sensitivity:\t\t\t', three_decimals[1])
print('  specificity:\t\t\t', three_decimals[2])
print('  positive likelihood:\t\t', three_decimals[3])
print('  negative likelihood:\t\t', three_decimals[4])
print('  false positive rate:\t\t', three_decimals[5])
print('  false negative rate:\t\t', three_decimals[6])
print('  positive predictive value:\t', three_decimals[7])
print('  negative predicitive value:\t', three_decimals[8])
print()

# %% Calculate weights from weights ratio:
# Set up class weighting to bias for sensitivty vs. specificity
# Higher values increase sensitivity at the cost of specificity
def calculate_class_weights(positive_class_weight_ratio):
    positive_weight = ( positive_class_weight_ratio /
                        (1 + positive_class_weight_ratio))

    negative_weight = 1 - positive_weight
    class_weights = {0: negative_weight, 1: positive_weight}
    return (class_weights)

### Create results folder if needed
# (Not used in this demo)
# OUTPUT_LOCATION = 'results'
# if not os.path.exists(OUTPUT_LOCATION):
#     os.makedirs(OUTPUT_LOCATION)

# %% Import data
url = ("https://raw.githubusercontent.com/MichaelAllen1966/wordpress_blog" +
      "/master/jupyter_notebooks/stroke.csv")
df_stroke = pd.read_csv(url)
feat_labels = list(df_stroke)[1:]
number_of_features = len(feat_labels)
X, y = df_stroke.iloc[:, 1:].values, df_stroke.iloc[:, 0].values

# Set different weights for pisitive and negative results in SVM is required
# This will adjust balance between sensitivity and specificity
# For equal weighting, set at 1
positive_class_weight_ratio = 1
class_weights = calculate_class_weights(positive_class_weight_ratio)

# Set up strtified k-fold
splits = 10
skf = StratifiedKFold(n_splits = splits)

```

```

skf.get_n_splits(X, y)

# %% Set up results dataframes
forest_results = np.zeros((splits, 10))
forest_importance = np.zeros((splits, number_of_features))
svm_results_linear = np.zeros((splits, 10))
svm_results_rbf = np.zeros((splits, 10))
lr_results = np.zeros((splits, 10))
nn_results = np.zeros((splits, 10))

# %% Loop through the k splits of training/test data
loop_count = 0

for train_index, test_index in skf.split(X, y):

    print ('Split', loop_count + 1, 'out of', splits)

    X_train, X_test = X[train_index], X[test_index]
    y_train, y_test = y[train_index], y[test_index]

    sc = StandardScaler() # new Standard Scalar object
    sc.fit(X_train)
    X_train_std = sc.transform(X_train)
    X_test_std = sc.transform(X_test)
    combined_results = pd.DataFrame()

    # %% Random forests
    forest = RandomForestClassifier(n_estimators=1000, n_jobs=-1,
                                   class_weight='balanced')

    forest.fit(X_train, y_train)
    forest_importance[loop_count, :] = forest.feature_importances_
    y_pred = forest.predict(X_test)
    test_results = pd.DataFrame(np.vstack((y_test, y_pred)).T)
    diagnostic_performance = (calculate_diagnostic_performance
                              (test_results.values))
    forest_results[loop_count, :] = diagnostic_performance
    combined_results['Forest'] = y_pred

    # %% SVM (Support Vector Machine) Linear
    svm = SVC(kernel='linear', C=1.0, class_weight=class_weights)
    svm.fit(X_train_std, y_train)
    y_pred = svm.predict(X_test_std)
    test_results = pd.DataFrame(np.vstack((y_test, y_pred)).T)
    diagnostic_performance = (calculate_diagnostic_performance
                              (test_results.values))
    svm_results_linear[loop_count, :] = diagnostic_performance
    combined_results['SVM_linear'] = y_pred

```

```

# %% SVM (Support Vector Machine) RBF
svm = SVC(kernel='rbf', C=1.0)
svm.fit(X_train_std, y_train)
y_pred = svm.predict(X_test_std)
test_results = pd.DataFrame(np.vstack((y_test, y_pred)).T)
diagnostic_performance = (calculate_diagnostic_performance
                          (test_results.values))
svm_results_rbf[loop_count, :] = diagnostic_performance
combined_results['SVM_rbf'] = y_pred

# %% Logistic Regression
lr = LogisticRegression(C=100, class_weight=class_weights)
lr.fit(X_train_std, y_train)
y_pred = lr.predict(X_test_std)
test_results = pd.DataFrame(np.vstack((y_test, y_pred)).T)
diagnostic_performance = (calculate_diagnostic_performance
                          (test_results.values))
lr_results[loop_count, :] = diagnostic_performance
combined_results['LR'] = y_pred

# %% Neural Network
clf = MLPClassifier(solver='lbfgs', alpha=1e-8, hidden_layer_sizes=(50, 5),
                   max_iter=100000, shuffle=True, learning_rate_init=0.001,
                   activation='relu', learning_rate='constant', tol=1e-7)
clf.fit(X_train_std, y_train)
y_pred = clf.predict(X_test_std)
test_results = pd.DataFrame(np.vstack((y_test, y_pred)).T)
diagnostic_performance = (calculate_diagnostic_performance
                          (test_results.values))
nn_results[loop_count, :] = diagnostic_performance
combined_results['NN'] = y_pred

# Increment loop count
loop_count += 1

# %% Transfer results to Pandas arrays
results_summary = pd.DataFrame()

results_column_names = (['accuracy', 'sensitivity',
                        'specificity',
                        'positive likelihood',
                        'negative likelihood',
                        'false positive rate',
                        'false negative rate',
                        'positive predictive value',
                        'negative predictive value',
                        'positive rate'])

```

```

forest_results_df = pd.DataFrame(forest_results)
forest_results_df.columns = results_column_names
forest_importance_df = pd.DataFrame(forest_importance)
forest_importance_df.columns = feat_labels
results_summary['Forest'] = forest_results_df.mean()

svm_results_lin_df = pd.DataFrame(svm_results_linear)
svm_results_lin_df.columns = results_column_names
results_summary['SVM_lin'] = svm_results_lin_df.mean()

svm_results_rbf_df = pd.DataFrame(svm_results_rbf)
svm_results_rbf_df.columns = results_column_names
results_summary['SVM_rbf'] = svm_results_rbf_df.mean()

lr_results_df = pd.DataFrame(lr_results)
lr_results_df.columns = results_column_names
results_summary['LR'] = lr_results_df.mean()

nn_results_df = pd.DataFrame(nn_results)
nn_results_df.columns = results_column_names
results_summary['Neural'] = nn_results_df.mean()

# %% Print summary results
print()
print('Results Summary:')
print(results_summary)

# %% Save files
# NOT USED IN THIS DEMO
# forest_results_df.to_csv('results/forest_results.csv')
# forest_importance_df.to_csv('results/forest_importance.csv')
# svm_results_lin_df.to_csv('results/sum_lin_results.csv')
# svm_results_rbf_df.to_csv('results/sum_rbf_results.csv')
# lr_results_df.to_csv('results/logistic_results.csv')
# nn_results_df.to_csv('results/neural_network_results.csv')
# results_summary.to_csv('results/results_summary.csv')

```

```

Split 1 out of 10
Split 2 out of 10
Split 3 out of 10
Split 4 out of 10
Split 5 out of 10
Split 6 out of 10
Split 7 out of 10
Split 8 out of 10
Split 9 out of 10
Split 10 out of 10

```

Results Summary:

	Forest	SVM_lin	SVM_rbf	LR	Neural
accuracy	0.851946	0.839995	0.843081	0.839610	0.801859
sensitivity	0.727978	0.767511	0.741951	0.753473	0.702353
specificity	0.905567	0.871350	0.886804	0.876865	0.844867
positive likelihood	8.799893	7.396559	7.384775	7.390298	4.909178
negative likelihood	0.297522	0.263269	0.287613	0.276478	0.349459
false positive rate	0.094433	0.128650	0.113196	0.123135	0.155133
false negative rate	0.272022	0.232489	0.258049	0.246527	0.297647
positive predictive value	0.775919	0.731363	0.747641	0.737093	0.669270
negative predictive value	0.887152	0.898619	0.890479	0.894471	0.869677
positive rate	0.285678	0.321505	0.302999	0.313414	0.320310