

Python for Healthcare Modelling

Michael Allen

March 28, 2018

Contents

1	Python basics	1
1.1	Introduction	2
1.2	Lists	3
1.2.1	Creating a list	3
1.2.2	Deleting a list	3
1.2.3	Returning and printing an element from a list	3
1.2.4	Taking slices of a list	3
1.2.5	Deleting or changing an item in a list	4
1.2.6	Appending an item to a list, and joining two lists	4
1.2.7	Copying a list	5
1.2.8	Inserting an element into a given position in a list	5
1.2.9	Sorting a list	5
1.2.10	Checking whether the list contains a given element	6
1.2.11	Reversing a list	6
1.2.12	Counting the number of elements in a list	6
1.2.13	Returning the index position of an element of a list	6
1.2.14	Removing an element of a list by its value	7
1.2.15	'Popping' an element from a list'	7
1.2.16	Iterating (stepping) through a list	7
1.2.17	Iterating through a list and changing element values	8
1.2.18	Counting the number of times an element exists in a list	8
1.2.19	Turning a sentence into a list of words	8
1.3	Nested lists	9
1.4	Tuples	10
1.4.1	Creating and adding to tuples	10
1.4.2	Converting between tuples and lists, and sorting tuples	10
1.5	Sets	12
1.5.1	Creating sets	12
1.5.2	Adding to a set and removing duplicates	12
1.5.3	Analysing the intersection between sets	12
1.6	Dictionaries	14
1.6.1	Creating and adding to dictionaries	14
1.6.2	Changing a dictionary entry	14
1.6.3	Deleting a dictionary entry	15
1.6.4	Iterating through a dictionary	15
1.6.5	Converting dictionaries to lists of keys and values	15
1.6.6	Using the get method to return a null value if a key does not exist	16
1.6.7	Intersection between dictionaries	16
1.6.8	Other dictionary methods	17
1.6.9	Ordered dictionaries	17
1.7	Accessing math functions through the math module	18
1.8	Variable types	24
1.9	Random numbers and integers	26
1.9.1	Importing the random module and setting a random seed	26
1.9.2	Random numbers	26
1.9.3	Generating random sequences	26

1.10	if, elif, else, while, and logical operators	28
1.10.1	if, else, elif	28
1.10.2	while statements	28
1.10.3	Logical operators	28
1.11	Loops and iterating	30
1.11.1	Breaking out of loops or continuing the loop without action	30
1.11.2	Using pass to replace active code in a loop	31
1.12	List comprehensions: one line loops	32
1.13	Else after while	33
1.14	try ... except (where code might fail)	34
1.15	Decimal places in output	35
1.16	Writing to and reading from files	36
1.16.1	File access modes	36
1.16.2	Creating a file with write	36
1.16.3	Adding to a file	36
1.16.4	Writing multiple lines	37
1.16.5	Reading a file	37
1.17	Functions	38
1.17.1	Defining a function	38
1.17.2	Adding help to a function with a docstring	39
1.17.3	Returning two or more results	39
1.18	Lambda functions (one line functions), and map/filter/reduce	41
1.18.1	Defining one-line lambda functions	41
1.18.2	Applying a lambda function to a list with map	41
1.18.3	Filtering a list with a lambda function	42
1.18.4	Reducing a list with a repeated lambda function call	42
1.18.5	Alternatives to map, filter and reduce	42
1.19	Accessing date and time, and timing code	43
1.19.1	Accessing date and time	43
1.19.2	Timing code	44
2	NumPy and Pandas	45
2.1	Numpy and Pandas	46
2.1.1	NumPy	46
2.1.2	Pandas	46
2.2	Numpy basics: building an array from lists, basic statistics, converting to booleans, refer- encing the array, and taking slicesi	47
2.2.1	Building an array from lists	47
2.2.2	Performing basic statistics on the array	47
2.2.3	Converting values to a boolean True/False	49
2.2.4	Referencing arrays and taking slices	49

Chapter 1

Python basics

1.1 Introduction

Hello

This book is a collection of code snippets that help me use Python for health services research, modelling and analysis. When learning something new I always work on a small code example to understand how something works, and to keep as a handy reference.

I will be looking at data handling, some statistics, data plotting, discrete event simulation, and machine learning.

I will be including use of pure Python as well as commonly used libraries such as NumPy, Pandas, Matplotlib, SciPy, SciKitLearn, TensorFlow and Simpy.

Everything described here is performed in Python 3, based on the Anaconda Scientific Python environment, available for free (for Windows, Mac or Linux). Follow installation instructions from the site <https://www.anaconda.com/download/>

Anaconda comes with its own environments for writing the code: Spyder or Jupyter Notebooks. Both are nice. Other Free and Open Source options are PyCharm for more functionality, Atom, or Visual Studio Code for a modern code text editor, or Vim for a more old-fashioned but very fast and lightweight code text editor.

Occasionally other free libraries may be installed. If so they will be described where appropriate.

I choose to do all my work in GNU/Linux, but everything should also work in Microsoft Windows or Mac OS.

If you are new to this I would recommend installing the Anaconda Scientific Python environment, and then look for 'Spyder' in your computer's application listing. That will open up an 'Integrated Development Environment' (IDE): a posh phrase that means a place where you can both write and run code.

Happy coding!

Michael

1.2 Lists

Lists, tuples, sets and dictionaries are four basic objects for storing data in Python. Later we'll look at libraries that allow for specialised handling and analysis of large or complex data sets.

Lists are a group of text or numbers (or more complex objects) held together in a given order. Lists are recognised by using square brackets, with members of the list separated by commas. Lists are mutable, that is their contents may be changed. Note that when Python numbers a list the first element is referenced by the index zero. Python is a 'zero-indexed' language.

Below is code demonstrating some common handling of lists.

1.2.1 Creating a list

Lists may mix text, numbers, or other Python objects. Here is code to generate and print a list of mixed text and numbers. Note that the text is in speech marks (single or double; both will work). If the speech marks were missing then Python would think the text represented a variable name (such as the name of another list).

```
my_list = ['Gandalf', 'Bilbo', 'Gimli', 10, 30, 40]
print (my_list)
```

OUT:

```
['Gandalf', 'Bilbo', 'Gimli', 10, 30, 40]
```

1.2.2 Deleting a list

The *del* command will delete a list (or any other variable).

```
my_list = ['Gandalf', 'Bilbo', 'Gimli']
del my_list
```

1.2.3 Returning and printing an element from a list

Python lists store their contents in sequence. A particular element may be referenced by its index number. Python indexes start at zero. Watch out for this this will almost certainly trip you up at some point.

The example below prints the second item in the list. The first item would be referenced with 0 (zero).

```
my_list = ['Gandalf', 'Bilbo', 'Gimli', 10, 30, 40]
print (my_list[1])
```

OUT:

```
Bilbo
```

Negative references may be used to refer to position from the end of the list. An index of -1 would be the last item (because -0 is the same as 0 and would be the first item!)

1.2.4 Taking slices of a list

Taking more than one element from a list is called a slice. The slice defines the starting point, and the point just after the end of the slice. The example below takes elements indexed 2 to 4 (but not 5)

```
my_list = [0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10]
slice = my_list[2:5]
print(slice)
```

OUT:

```
[2, 3, 4]
```

Slices can take every *n* elements, the example below takes every second element from index 1 up to, but not including, index 9 (remembering that the first element is index 0):

```
y_list = [0,1,2,3,4,5,6,7,8,9,10]
slice = my_list[2:9:2]
print(slice)
```

OUT:

```
[2, 4, 6, 8]
```

Reverse slices may also be taken by defining the end point before the beginning point and using a negative step:

```
my_list = [0,1,2,3,4,5,6,7,8,9,10]
slice = my_list[9:2:-2]
print(slice)
```

OUT:

```
[9, 7, 5, 3]
```

1.2.5 Deleting or changing an item in a list

Unlike tuples (which are very similar to lists), elements in a list may be deleted or changed.

Here the first element in the list is deleted:

```
my_list = ['Gandalf','Bilbo','Gimli',10,30,40]
del my_list[0]
print (my_list)
```

OUT:

```
['Bilbo', 'Gimli', 10, 30, 40]
```

Here the second element of the list is replaced:

```
my_list = ['Gandalf','Bilbo','Gimli',10,30,40]
my_list[0] = 'Boromir'
print (my_list)
```

OUT:

```
['Boromir', 'Bilbo', 'Gimli', 10, 30, 40]
```

1.2.6 Appending an item to a list, and joining two lists

Individual elements may be added to a list with *append*.

```
my_list = ['Gandalf','Bilbo','Gimli',10,30,40]
my_list.append('Elrond')
print (my_list)
```

OUT:

```
['Gandalf', 'Bilbo', 'Gimli', 10, 30, 40, 'Elrond']
```

Two lists may be joined with a simple +

```
my_list = ['Gandalf','Bilbo','Gimli']
my_second_list = ['Elrond','Boromir']
my_combined_list = my_list + my_second_list
print (my_combined_list)
```


OUT:

```
['Gandalf', 'Bilbo', 'Gimli', 'Elrond', 'Boromir']
```

Or the *extend* method may be used to add a list to an existing list.

```
my_list = ['Gandalf', 'Bilbo', 'Gimli']
my_second_list = ['Elrond', 'Boromir']
my_list.extend(my_second_list)
print (my_list)
```

OUT:

```
['Gandalf', 'Bilbo', 'Gimli', 'Elrond', 'Boromir']
```

1.2.7 Copying a list

If we try to copy a list by saying `mycopy mylist` we do not generate a new copy. Instead we generate an alias, another name that refers to the original list. Any changes to `mycopy` will change `mylist`.

There are two ways we can make a separate copy of a list, where changes to the new copy will not change the original list:

```
my_list = ['Gandalf', 'Bilbo', 'Gimli']
my_copy = my_list.copy()
# Or
my_copy = my_list[:]
```

1.2.8 Inserting an element into a given position in a list

Inserting an element into a given position in a list

An element may be inserted into a list with the *insert* method. Here we specify that the new element will be inserted at index 2 (the third element in the list, remembering that the first element is index 0). In a later post we will look at methods specifically for maintaining a sorted list, but for now let us simply inset an element at a given position in the list.

```
my_list = ['Gandalf', 'Bilbo', 'Gimli']
my_list.insert(2, 'Boromir')
print (my_list)
```

OUT:

```
['Gandalf', 'Bilbo', 'Boromir', 'Gimli']
```

1.2.9 Sorting a list

The *sort* method can work on a list of text or a list of numbers (but not a mixed list text and numbers). Sort order may be reversed if wanted.

A normal sort:

```
x = [10, 5, 24, 5, 8]
x.sort()
print (x)
```

OUT:

```
[5, 5, 8, 10, 24]
```

A reverse sort:

```
x = [10, 5, 24, 5, 8]
x.sort(reverse=True)
print (x)
```

```
OUT:
[24, 10, 8, 5, 5]
```

1.2.10 Checking whether the list contains a given element

The *in* command checks whether a particular element exist in a list. It returns a 'Boolean' True or False.

```
my_list = ['Gandalf', 'Bilbo', 'Gimli', 'Elrond', 'Boromir']
is_in_list = 'Elrond' in my_list
print (is_in_list)
is_in_list = 'Bob' in my_list
print (is_in_list)
```

```
OUT:
True
False
```

1.2.11 Reversing a list

The method to reverse a list looks a little odd. It is actually creating a 'slice' of a list that steps back from the end of list.

```
my_list = ['Gandalf', 'Bilbo', 'Gimli', 'Elrond', 'Boromir']
my_list = my_list[::-1]
print (my_list)
```

```
OUT:
['Boromir', 'Elrond', 'Gimli', 'Bilbo', 'Gandalf']
```

1.2.12 Counting the number of elements in a list

Use the *len* function to return the number of elements in a list.

```
my_list = ['Gandalf', 'Bilbo', 'Gimli', 'Elrond', 'Boromir']
print (len(my_list))
```

```
OUT:
5
```

1.2.13 Returning the index position of an element of a list

The list *index* will return the position of the first element of the list matching the argument given.

```
my_list = ['Gandalf', 'Bilbo', 'Gimli', 'Elrond', 'Boromir']
index_pos = my_list.index('Elrond')
print (index_pos)
```

```
OUT:
3
```

If the element is not in the list then the code will error. Later we will look at some different ways of coping with this type of error. For now let us introduce an if/then/else statement (and we'll cover those in more detail later as well).

```
my_list = ['Gandalf','Bilbo','Gimli','Elrond','Boromir']
test_element = 'Elrond'
is_in_list = test_element in my_list
if is_in_list: # this is the same as saying if is_in_list == True
    print (test_element,'is in list at position', my_list.index(test_element))
else:
    print (test_element, 'is not in list')
```

OUT:

Elrond is in list at position 3

1.2.14 Removing an element of a list by its value

An element may be removed from a list by using its value rather than index. Be aware that you might want to put in a method (like the if statement above) to check the value is present before removing it. But here is a simple remove (this will remove the first instance of that particular value:

```
my_list = ['Gandalf','Bilbo','Gimli','Elrond','Boromir']
my_list.remove('Bilbo')
print (my_list)
```

OUT:

['Gandalf', 'Gimli', 'Elrond', 'Boromir']

1.2.15 'Popping' an element from a list

Popping an element from a list means to retrieve an element from a list, removing it from the list at the same time. If no element number is given then the last element of the list will be popped.

```
my_list = ['Gandalf','Bilbo','Gimli','Elrond','Boromir']
print ('Original list:', my_list)
x = my_list.pop()
print (x, 'was popped')
print ('Remaining list:',my_list)
x = my_list.pop(1)
print (x, 'was popped')
print ('Remaining list:', my_list)
```

OUT:

Original list: ['Gandalf', 'Bilbo', 'Gimli', 'Elrond', 'Boromir']
 Boromir was popped
 Remaining list: ['Gandalf', 'Bilbo', 'Gimli', 'Elrond']
 Bilbo was popped
 Remaining list: ['Gandalf', 'Gimli', 'Elrond']

1.2.16 Iterating (stepping) through a list

Lists are 'iterable', that is they can be stepped through. The example below prints one element at a time. We use the variable name 'x' here, but any name may be used.

```
my_list = [1,2,3,4,5,6,7,8,9]
for x in my_list:
    print(x, end=' ') # end command replaces a newline with a space
```

OUT:

```
1 2 3 4 5 6 7 8 9
```

1.2.17 Iterating through a list and changing element values

Iterating through a list and changing the value in the original list is a little more complicated. Below the range command creates a range of numbers from zero up to the index of the last element of the list. If there were 4 elements, for example, then the range function would produce '0, 1, 2, 3'. This allows us to iterate through the index numbers for the list (the position of each element). Using the index number we can change the value of the element in that position of the list.

```
my_list = [1,2,3,4,5,6,7,8,9]
for index in range(len(my_list)):
    my_list[index] = my_list[index]**2
print ('\nMy list after iterating and mutating:')
print (my_list)
```

OUT:

```
My list after iterating and mutating:
[1, 4, 9, 16, 25, 36, 49, 64, 81]
```

1.2.18 Counting the number of times an element exists in a list

The *count* method counts the number of times a value occurs in a list. If the value does not exist a zero is returned.

```
x =[1, 2, 3, 4, 1, 4, 1]
test_value = 1
print ('Count the number of times 1 occurs in',x)
print (x.count(test_value),'\n')
```

OUT:

```
Count the number of times 1 occurs in [1, 2, 3, 4, 1, 4, 1]
3
```

1.2.19 Turning a sentence into a list of words

The *split* method allows a long string, such as a sentence, to be separated into individual words. If no separator is defined any white space (such as a space or comma) will be used to divide words.

```
my_sentence ="Turn this sentence into a list"
words = my_sentence.split()
print (words)
```

OUT:

```
['Turn', 'this', 'sentence', 'into', 'a', 'list']
```

A separator may be specified, such as dividing sentences at commas:

```
my_sentence = 'Hello there, my name is Bilbo'
divided_sentence = my_sentence.split(sep=',')
print(divided_sentence)
```

OUT:

```
['Hello there', ' my name is Bilbo']
```

1.3 Nested lists

So far we have looked at lists which contain a simple series of numbers, text, or a mixture of numbers and texts (Python lists can also hold any Python object, but in Healthcare modelling we are usually dealing with numbers or text in a list).

It is possible though to build nested lists. In the example below we generate a list manually, with each nested list on a separate line. This separation by line is just to make it easier to see; it is not needed in Python, but thought to layout of code is important if other people will be looking at your code.

```
my_list = [[1,2,3],
           [4,5,6],
           [7,8,9]]
print (my_list)
```

OUT:

```
[[1, 2, 3], [4, 5, 6], [7, 8, 9]]
```

We can then refer to items by using two reference indices. The first refers to the nested list block (so the first id [1, 2, 3] and the second refers to the position within that block. Remember that Python is zero indexed so the first element of the first list is 0[0].

```
print (my_list[1][1])
```

OUT:

```
5
```

Or we can use other variables to refer to the position:

```
x = 2
y = 0
print (my_list[x][y])
```

OUT:

```
7
```

More complex structures can be built up with further nesting of lists to give multi-dimensional lists.

WARNING: Handling large arrays of data this way is possible but slow. For modelling we are much better off using two libraries dedicated to fast handling of large data sets: NumPy and Pandas. We will be covering those libraries soon.

1.4 Tuples

1.4.1 Creating and adding to tuples

Tuples are like lists in many ways, apart from they cannot be changed - they are immutable. Tuples are defined using curved brackets (unlike square brackets for lists). Tuples may be returned, or be required, by various functions in Python, and tuples may be chosen where immutability would be an advantage (for example, to help prevent accidental changes).

```
my_tuple = ('Hobbit', 'Elf', 'Ork', 'Dwarf')
print (my_tuple[1])
```

OUT:
Elf

It is possible to add to a tuple. Note that if adding a single item an additional comma is used to indicate to Python that the variable being added is a tuple.

```
my_tuple = my_tuple + ('Man',)
my_tuple += ('Wizard', 'Balrog') # Note that the += is short hand to add something to itself
print (my_tuple)
```

OUT:
('Hobbit', 'Elf', 'Ork', 'Dwarf', 'Man', 'Wizard', 'Balrog')

It is not possible to change or delete an item in a tuple. To change or delete a tuple a new tuple must be built (but if this is going to happen then a list would be a better choice).

```
my_new_tuple = my_tuple[0:2] + ('Goblin',) + my_tuple[4:]
print (my_new_tuple)
```

OUT:
('Hobbit', 'Elf', 'Goblin', 'Man', 'Wizard', 'Balrog')

1.4.2 Converting between tuples and lists, and sorting tuples

A tuple may be turned into a list. We can recognise that it is a list by the square brackets.

```
my_list = list(my_tuple)
print (my_list)
```

OUT:
['Hobbit', 'Elf', 'Ork', 'Dwarf', 'Man', 'Wizard', 'Balrog']

A list may also be converted into a tuple.

```
my_list.sort()
my_new_tuple = tuple(my_list)
print (my_new_tuple)
```

OUT:
my_list.sort()

```
my_new_tuple = tuple(my_list)
```

```
print (my_new_tuple)
```

('Balrog', 'Dwarf', 'Elf', 'Hobbit', 'Man', 'Ork', 'Wizard')

In the above example we sorted a list and converted it to a tuple. Tuples cannot be changed (apart from being added to), so it is not possible to directly sort a tuple. The sorted command will act on a tuple to

sort it, but returns a list.

```
print (sorted(my_tuple))
```

OUT:

```
['Balrog', 'Dwarf', 'Elf', 'Hobbit', 'Man', 'Ork', 'Wizard']
```

The tuple may be sorted by converting back to a list in a single step (but think of using lists, rather than tuples, is sorting will be common.

```
my_tuple = tuple(sorted(my_tuple))  
print (my_tuple)
```

OUT:

```
('Balrog', 'Dwarf', 'Elf', 'Hobbit', 'Man', 'Ork', 'Wizard')
```

1.5 Sets

Sets are unordered collections of unique values. Common uses include membership testing, finding overlaps and differences between sets, and removing duplicates from a sequence.

Sets, like dictionaries, are defined using curly brackets.

1.5.1 Creating sets

Creating a set directly:

```
my_set = {'a','e','i','o','u'}  
print (my_set)
```

OUT:

```
{'a', 'i', 'u', 'o', 'e'}
```

Creating a set from a list:

```
my_list = ['a','e','i','o','u']  
my_set = set(my_list)  
print (my_set)
```

OUT:

```
{'a', 'i', 'u', 'o', 'e'}
```

Creating an empty set. This cannot be done simply with

```
my_set = set()  
print (my_set)
```

1.5.2 Adding to a set and removing duplicates

Adding one element to a set:

```
my_set.add('y')  
print (my_set)
```

OUT:

```
{'a', 'y', 'i', 'u', 'o', 'e'}
```

Adding multiple elements to a set (note that the set has no duplication of entries already present; sets automatically remove duplicate items):

```
my_set.add('y')  
print (my_set)
```

OUT:

```
{'a', 'y', 'i', 'u', 'o', 'e'}
```

1.5.3 Analysing the intersection between sets

Sets may be used to explore the intersection between sets. There are usually two forms of syntax which may be used, both of which are given in the examples below.

```
set1 = {'a','b','c','d','e','f'}  
set2 = {'a','e','i','o','u'}
```

The difference between two sets:


```
print ('difference')
print (set1.difference(set2))
print (set1 - set2)
```

```
OUT:
difference
{'c', 'f', 'b', 'd'}
{'c', 'f', 'b', 'd'}
```

The intersection between two sets:

```
print ('Intersection')
print (set1.intersection(set2))
print (set1 & set2)
```

```
OUT:
Intersection
{'a', 'e'}
{'a', 'e'}
```

Is a subset? Is set 2 wholly included in set 1?

```
print ('Is subset?')
print (set1.issubset(set2))
print (set1 <= set2)
```

```
OUT:
Is subset?
False
False
```

Is a superset? Does set 1 wholly include set 2?

```
print ('Is superset?')
print (set1.issuperset(set2))
print (set1 >= set2)
```

```
OUT:
Is superset?
False
False
```

The union between two sets:

```
print ('Union')
print (set1.union(set2))
print (set1 | set2)
```

```
OUT:
Union
{'b', 'o', 'a', 'f', 'c', 'i', 'u', 'd', 'e'}
{'b', 'o', 'a', 'f', 'c', 'i', 'u', 'd', 'e'}
```

Symmetric difference: in either but not both (equivalent to xor)

```
print ('Symmetric difference')
print (set1.symmetric_difference(set2))
print (set1 ^ set2)
```

```
OUT:
Symmetric difference
{'c', 'i', 'b', 'u', 'f', 'o', 'd'}
{'c', 'i', 'b', 'u', 'f', 'o', 'd'}
```

1.6 Dictionaries

Dictionaries store objects such as text or numbers (or other Python objects). We saw that lists stored elements in sequence, and that each element can be referred to by an index number, which is the position of that element in a list. The contents of dictionaries are accessed by a unique *key* which may be a number or may be text.

Dictionaries take the form of `key1:value1, key2:value2`

The *value* may be a single number or word, or may be another Python object such as a list or a tuple.

Dictionaries allow fast random-access to data.

1.6.1 Creating and adding to dictionaries

Dictionaries may be created with content, and are referred to by their *key*:

```
fellowship = {'Man': 2, 'Hobbit':4, 'Wizard':1, 'Elf':1, 'Dwarf':1}
print (fellowship['Man'])
```

OUT:
2

Dictionaries may also be created empty. New content may be added to an empty or an existing dictionary.

```
trilogy = {}
trilogy['Part 1'] = 'The fellowship of the ring'
trilogy['Part 2'] = 'The Two Towers'
trilogy['Part 3'] = 'The return of the king'
print (trilogy['Part 2'])
```

OUT:
The Two Towers

Dictionaries may contain a mixture of data types.

```
lord_of_the_rings = {'books':3,
                    'author':'JRR Tolkien',
                    'main_characters':['Frodo', 'Sam', 'Gandalf']}
print (lord_of_the_rings['main_characters'])
```

OUT:
['Frodo', 'Sam', 'Gandalf']

1.6.2 Changing a dictionary entry

Dictionary elements are over-written if a key already exists.

```
trilogy['Part 2'] = 'The Two Towers'
print (trilogy['Part 2'])
```

OUT:
The Two Towers

Dictionary items may also be updated using another dictionary:

```
entry_update = {'Part 1':'The Fellowship of the Ring',
                'Part 3':'The Return of the King'}
trilogy.update(entry_update)
print (trilogy)
```

OUT:

```
{'Part 1': 'The Fellowship of the Ring', 'Part 2': 'The Two Towers',
'Part 3': 'The Return of the King'}
```

1.6.3 Deleting a dictionary entry

Dictionary entries may be deleted, by reference to their key, with the *del* command. If the key does not exist an error will be caused, so in the example below we check that the entry exists before trying to delete it.

```
entry_to_delete = 'Part 2'
if entry_to_delete in trilogy:
    del trilogy[entry_to_delete]
    print (trilogy)
else:
    print('That key does not exist')
```

OUT:

```
{'Part 1': 'The Fellowship of the Ring', 'Part 3': 'The Return of the King'}
```

1.6.4 Iterating through a dictionary

There are two main methods to iterate through a dictionary. The first retrieves just the key, and the value may then be retrieved from that key. The second method retrieves the key and value together.

```
lord_of_the_rings = {'books':3,
                    'author':'JRR Tolkien',
                    'main_characters':['Frodo','Sam','Gandalf']}

# Iterating method 1
for key in lord_of_the_rings:
    print (key, '-', lord_of_the_rings[key])

# Iterating method 2
print()
for key, value in lord_of_the_rings.items():
    print (key, '-', value)
```

OUT:

```
books - 3
author - JRR Tolkien
main_characters - ['Frodo', 'Sam', 'Gandalf']
```

```
books - 3
author - JRR Tolkien
main_characters - ['Frodo', 'Sam', 'Gandalf']
```

1.6.5 Converting dictionaries to lists of keys and values

Dictionary keys and values may be accessed separately. We can also return a list of tuples of keys and values with the *items* method.

```
lord_of_the_rings = {'books':3,
                    'author':'JRR Tolkien',
                    'main_characters':['Frodo','Sam','Gandalf']}

# print the keys:
```

```
print ('\nKeys:') # \n creates empty line before text
print (list(lord_of_the_rings.keys()))
```

```
# print the values
print ('\nValues:')
print (list(lord_of_the_rings.values()))
```

```
# print keys and values
print ('\nKeys and Values:')
print (list(lord_of_the_rings.items()))
```

OUT:

```
Keys:
['books', 'author', 'main_characters']
```

Values:

```
[3, 'JRR Tolkien', ['Frodo', 'Sam', 'Gandalf']]
```

Keys and Values:

```
[('books', 3), ('author', 'JRR Tolkien'), ('main_characters', ['Frodo', 'Sam', 'Gandalf'])]
```

1.6.6 Using the get method to return a null value if a key does not exist

Using the usual method of referring to a dictionary element, an error will be returned if the key used does not exist. An 'if key in dictionary' statement could be used to check the key exists. Or the get method will return a null value if the key does not exist. The get method allows for a default value to be returned if a key is not found.

```
print (lord_of_the_rings.get('author'))
print (lord_of_the_rings.get('elves'))
print (lord_of_the_rings.get('books','Key not found - sorry'))
print (lord_of_the_rings.get('dwarfs','Key not found - sorry'))
```

JRR Tolkien

None

3

Key not found - sorry

1.6.7 Intersection between dictionaries

Like sets the intersection between dictionaries may be evaluated. See the entry on sets for more intersection methods.

```
a = {'x':1,'y':2,'z':3,'zz':2}
b = {'x':10,'w':11,'y':2,'yy':2}
# Show keys in common
print('Keys in common: ',a.keys() & b.keys())
# Show keys in a not in b
print('Keys in a but not b:',a.keys()-b.keys())
# Show keys+values in common (items() return key+value, so both need to be the same)
print('Keys and values in common:',a.items() & b.items())
```

OUT:

```
Keys in common: {'x', 'y'}
```

```
Keys in a but not b: {'z', 'zz'}
```

```
Keys and values in common: {('y', 2)}
```

1.6.8 Other dictionary methods

`dictionary.clear()` empties a dictionary

`dictionary.pop(key)` will return an item and remove it from the dictionary

1.6.9 Ordered dictionaries

Like sets, dictionaries do not usually keep the order of entries. If it is useful to keep the order of entry in a dictionary then the `OrderedDict` object is useful. This needs an import statement before it is used (all import statements are usually kept together at the start of code).

```
from collections import OrderedDict
d=OrderedDict()
```

```
d['Gandalf']=1
d['Bilbo']=2
d['Frodo']=3
d['Sam']=4
```

```
for key in d:
    print(key,d[key])
```

OUT:

```
Gandalf 1
Bilbo 2
Frodo 3
Sam 4
```

1.7 Accessing math functions through the math module

Here we are going to use the math module as an introduction to using modules. The math module contains a range of useful mathematical functions that are not built into Python directly. So let's go ahead and start by importing the module. Module imports are usually performed at the start of a programme.

```
import math
```

When this type of import is used Python loads up a link to the module so that module functions and variables may be used. Here for example we access the value of pi through the module.

```
print (math.pi)
```

OUT:

```
3.141592653589793
```

Another way of accessing module contents is to directly load up a function or a variable into Python. When we do this we no longer need to use the module name after the import. This method is not generally recommended as it can lead to conflicts of names, and is not so clear where that function or variable comes from. But here is how it is done.

```
from math import pi
print (pi)
```

OUT:

```
3.141592653589793
```

Multiple methods and variables may be loaded at the same time in this way.

```
from math import pi, tau, log10
print (tau)
print (log10(100))
```

OUT:

```
6.283185307179586
2.0
```

But usually it is better practice to keep using the library name in the code.

```
print (math.log10(100))
```

OUT:

```
2.0
```

To access help on any Python function use the help command in a Python interpreter.

```
help (math.log10)
```

Help on built-in function log10 in module math:

```
log10(...)
    log10(x)
```

Return the base 10 logarithm of x.

To find out all the methods in a module, and how to use those methods we can simply type help (module.name) into the Python interpreter. The module must first have been imported, as we did for math above.

```
help (math)
```

OUT:

```
Help on module math:
```

NAME

math

MODULE REFERENCE

<https://docs.python.org/3.6/library/math>

The following documentation is automatically generated from the Python source files. It may be incomplete, incorrect or include features that are considered implementation detail and may vary between Python implementations. When in doubt, consult the module reference at the location listed above.

DESCRIPTION

This module is always available. It provides access to the mathematical functions defined by the C standard.

FUNCTIONS

`acos(...)`
`acos(x)`

Return the arc cosine (measured in radians) of x.

`acosh(...)`
`acosh(x)`

Return the inverse hyperbolic cosine of x.

`asin(...)`
`asin(x)`

Return the arc sine (measured in radians) of x.

`asinh(...)`
`asinh(x)`

Return the inverse hyperbolic sine of x.

`atan(...)`
`atan(x)`

Return the arc tangent (measured in radians) of x.

`atan2(...)`
`atan2(y, x)`

Return the arc tangent (measured in radians) of y/x.
Unlike `atan(y/x)`, the signs of both x and y are considered.

`atanh(...)`
`atanh(x)`

Return the inverse hyperbolic tangent of x.

`ceil(...)`
`ceil(x)`

Return the ceiling of x as an Integral.
This is the smallest integer $\geq x$.

```
copysign(...)  
    copysign(x, y)
```

Return a float with the magnitude (absolute value) of *x* but the sign of *y*. On platforms that support signed zeros, `copysign(1.0, -0.0)` returns `-1.0`.

```
cos(...)  
    cos(x)
```

Return the cosine of *x* (measured in radians).

```
cosh(...)  
    cosh(x)
```

Return the hyperbolic cosine of *x*.

```
degrees(...)  
    degrees(x)
```

Convert angle *x* from radians to degrees.

```
erf(...)  
    erf(x)
```

Error function at *x*.

```
erfc(...)  
    erfc(x)
```

Complementary error function at *x*.

```
exp(...)  
    exp(x)
```

Return *e* raised to the power of *x*.

```
expm1(...)  
    expm1(x)
```

Return `exp(x)-1`.

This function avoids the loss of precision involved in the direct evaluation of `exp(x)-1` for small *x*.

```
fabs(...)  
    fabs(x)
```

Return the absolute value of the float *x*.

```
factorial(...)  
    factorial(x) -> Integral
```

Find *x*!. Raise a `ValueError` if *x* is negative or non-integral.

```
floor(...)  
    floor(x)
```


Return the floor of x as an Integral.
This is the largest integer $\leq x$.

```
fmod(...)
fmod(x, y)
```

Return $\text{fmod}(x, y)$, according to platform C. $x \% y$ may differ.

```
frexp(...)
frexp(x)
```

Return the mantissa and exponent of x , as pair (m, e) .
 m is a float and e is an int, such that $x = m * 2.**e$.
If x is 0, m and e are both 0. Else $0.5 \leq \text{abs}(m) < 1.0$.

```
fsum(...)
fsum(iterable)
```

Return an accurate floating point sum of values in the iterable.
Assumes IEEE-754 floating point arithmetic.

```
gamma(...)
gamma(x)
```

Gamma function at x .

```
gcd(...)
gcd(x, y) -> int
greatest common divisor of x and y
```

```
hypot(...)
hypot(x, y)
```

Return the Euclidean distance, $\text{sqrt}(x*x + y*y)$.

```
isclose(...)
isclose(a, b, *, rel_tol=1e-09, abs_tol=0.0) -> bool
```

Determine whether two floating point numbers are close in value.

```
rel_tol
    maximum difference for being considered "close", relative to the
    magnitude of the input values
abs_tol
    maximum difference for being considered "close", regardless of the
    magnitude of the input values
```

Return True if a is close in value to b , and False otherwise.

For the values to be considered close, the difference between them must be smaller than at least one of the tolerances.

$-\text{inf}$, inf and NaN behave similarly to the IEEE 754 Standard. That is, NaN is not close to anything, even itself. inf and $-\text{inf}$ are only close to themselves.

```
isfinite(...)
isfinite(x) -> bool
```

Return True if x is neither an infinity nor a NaN, and False otherwise.

```
isinf(...)
    isinf(x) -> bool
```

Return True if x is a positive or negative infinity, and False otherwise.

```
isnan(...)
    isnan(x) -> bool
```

Return True if x is a NaN (not a number), and False otherwise.

```
ldexp(...)
    ldexp(x, i)
```

Return $x * (2^{**i})$.

```
lgamma(...)
    lgamma(x)
```

Natural logarithm of absolute value of Gamma function at x.

```
log(...)
    log(x[, base])
```

Return the logarithm of x to the given base.

If the base not specified, returns the natural logarithm (base e) of x.

```
log10(...)
    log10(x)
```

Return the base 10 logarithm of x.

```
log1p(...)
    log1p(x)
```

Return the natural logarithm of 1+x (base e).

The result is computed in a way which is accurate for x near zero.

```
log2(...)
    log2(x)
```

Return the base 2 logarithm of x.

```
modf(...)
    modf(x)
```

Return the fractional and integer parts of x. Both results carry the sign of x and are floats.

```
pow(...)
    pow(x, y)
```

Return x^{**y} (x to the power of y).

```
radians(...)
    radians(x)
```

Convert angle x from degrees to radians.

```
sin(...)
sin(x)
```

Return the sine of x (measured in radians).

```
sinh(...)
sinh(x)
```

Return the hyperbolic sine of x .

```
sqrt(...)
sqrt(x)
```

Return the square root of x .

```
tan(...)
tan(x)
```

Return the tangent of x (measured in radians).

```
tanh(...)
tanh(x)
```

Return the hyperbolic tangent of x .

```
trunc(...)
trunc(x:Real) -> Integral
```

Truncates x to the nearest Integral toward 0. Uses the `__trunc__` magic method.

DATA

```
e = 2.718281828459045
inf = inf
nan = nan
pi = 3.141592653589793
tau = 6.283185307179586
```

FILE

```
/home/michael/anaconda3/lib/python3.6/lib-dynload/math.cpython-36m-x86_64-linux-gnu.so
```

So now, for example, we know that to take a square root of a number we can use the `math` module, and use the `sqrt()` function, or use the `pow()` function which can do any power or root.

```
print (math.sqrt(4))
print (math.pow(4,0.5))
```

OUT:

```
2.0
2.0
```

In Python you might read about packages as well as modules. The two names are sometimes used interchangeably, but strictly a package is a collection of modules.

1.8 Variable types

Python is a dynamic language where variable type (e.g. whether a variable is a string or an integer) is not fixed. Sometime though you might like to force a particular type, or convert between types (most common where numbers may be contained within a string).

The most common types of variables in python are integers, float (nearly all types of numbers that are not integers) and strings.

Within python code the function *type* will show what variable type a string is.

```
x = 11 # This is an integer
print (type(x))
```

```
OUT:
<class 'int'>
```

Outside of the code, within the interpreter requesting help on a variable name will give the help for its variable type.

```
help (x)
```

```
OUT:
Help on int object:
```

```
class int(object)
|  int(x=0) -> integer
|  int(x, base=10) -> integer
|
|  Convert a number or string to an integer, or return 0 if no arguments
|  are given.  If x is a number, return x.__int__().  For floating point
|  numbers, this truncates towards zero.
|
|  If x is not a number or if base is given, then x must be a string,
|  bytes, or bytearray instance representing an integer literal in the
|  given base.  The literal can be preceded by '+' or '-' and be surrounded
|  by whitespace.  The base defaults to 10.  Valid bases are 0 and 2-36.
|  Base 0 means to interpret the base from the string as an integer literal.
|  >>> int('0b100', base=0)
|  4
|
|  Methods defined here:
```

```
...
```

Though x is an integer, if we divide it by 2, Python will dynamically change its type to float. This may be common sense to some, but may be irritating to people who are used to coding languages where variables have 'static' types (this particular behaviour is also different to Python2 where the normal behaviours is for integers to remain as integers).

```
x = x/2
print (x)
print (type(x))
```

```
OUT:
5.5
<class 'float'>
```

Python is trying to be helpful, but this dynamic behaviour of variables can sometimes need a little debugging. In our example here, to keep our answer as an integer (if that is what we needed to do) we need to specify that:

```
x = 11
x = int(x/2)
print (x)
print (type(x))
```

OUT:

```
5
<class 'int'>
```

We can convert between types, such as in the follow examples. Numbers may exist as text, particularly in some file imports.

```
x = '10'
print ('Look what happens when we try to multiply a string')
print (x *3)
x = float (x)
print ('Now let us multiply the converted variable')
print (x*3)
```

OUT:

```
Look what happens when we try to multiply a string
101010
Now let us multiply the converted variable
30.0
```

1.9 Random numbers and integers

Here we look at the standard Python random number generator. It uses a *Mersenne Twister*, one of the mostly commonly-used random number generators. The generator can generate random integers, random sequences, and random numbers according to a number of different distributions.

1.9.1 Importing the random module and setting a random seed

```
import random
```

Random numbers can be repeatable in the sequence they are generated if a 'seed' is given. If no seed is given the random number sequence generated each time will be different.

```
random.seed() # This will reset the random seed so that the random number sequence will be different
random.seed(10) # giving a number will lead to a repeatable random number sequence each time
```

1.9.2 Random numbers

Random integers between (and including) a minimum and maximum:

```
i
print (random.randint(0,5))
```

```
OUT:
4
```

Return a random number between 0 and 1:

```
print (random.random())
```

```
OUT:
0.032585065282054626
```

Return a number (floating, not integer) between a & b:

```
print (random.uniform(0,5))
```

```
OUT:
2.412808372754279
```

Select from normal distribution (with mu, sigma):

```
print (random.normalvariate(10,3))
```

```
OUT:
5.3538059684648855
```

Other distributions (see `help(random)` for more info after importing random module):

Lognormal, Exponential, Beta, Gaussian, Gamma, Weibul

1.9.3 Generating random sequences

The random library may also be used to shuffle a list:

```
deck = ['ace', 'two', 'three', 'four']
random.shuffle(deck)
print (deck)
```

```
OUT:
['three', 'ace', 'two', 'four']
```

Sampling without replacement:

```
sample = random.sample([10, 20, 30, 40, 50], k=4) # k is the number of samples to select
print (sample)
```

OUT:

```
[20, 30, 10, 50]
```

Sampling with replacement:

```
pick_from = ['red', 'black', 'green']
sample = random.choices(pick_from, k=10)
print(sample)
```

OUT:

```
['black', 'red', 'black', 'green', 'red', 'red', 'black', 'red', 'black', 'green']
```

Sampling with replacement and weighted sampling:

```
pick_from = ['red', 'black', 'green']
pick_weights = [18, 18, 2]
sample = random.choices(pick_from, pick_weights, k=10)
print(sample)
```

OUT:

```
['black', 'red', 'red', 'red', 'black', 'red', 'red', 'black', 'red', 'black']
```

1.10 if, elif, else, while, and logical operators

Like many programming languages, Python enables code to be run based on testing whether a condition is true. Python uses indented blocks to run code according to these tests.

1.10.1 if, else, elif

emphif statements run a block of code if a particular condition is true. emphelif or 'else if' statements can subsequently run to test more conditions, but these run only if none of the previous emphif or emphelif statements were true. emphelse statements may be used to run code if none of the previous emphif or emphelif were true.

emphif statements may be run alone, with emphelse statements or with emphelif statements. emphelse and emphelif statements require a previous emphif statement.

In the following example we use the input command to ask a user for a password and test against known value.

Here we use just one elif statement, but multiple statements could be used.

Note that the tests of equality uses double = signs

```
password = input ("Password? ") # get user to enter password
if password == "secret":
    print ("Well done")
    print ("You")
elif password=="Secret":
    print ("Close!")
else:
    print("Noodle brain")
```

OUT:

```
Password? Secret
Close!
```

1.10.2 while statements

emphwhile statements may be used to repeat some actions until a condition is no longer true. For example:

```
x = 0
while x <5:
    x +=1 # This is shorthand for x = x + 1
    print (x)
```

OUT:

```
1
2
3
4
5
```

1.10.3 Logical operators

Logical operators

The following are commonly used logical operators:

== Test of identity

!= Not equal to

>greater than

<less than

>= equal to or greater than

<= less than or equal to

in test whether element is in list/tuple/dictionary

not in test whether element is not in list/tuple/dictionary

and test multiple conditions are true

or test one of alternative conditions are true

any test whether all elements are true

all test whether all elements are true

When Python test conditions they are evaluated as either True or False. These values may also be used directly in tests:

```
x = 10
```

```
y = 20
```

```
z = 30
```

```
# using and/or
```

```
print (x>15)
```

```
print (x>15 and y>15 and z>15)
```

```
print (x>15 or y>15 or z>15)
```

```
print ()
```

```
# Using any and all
```

```
print (any([x>20, y>20, z>20]))
```

```
test_array = [x>20, y>20, z>20]
```

```
print (test_array)
```

```
print (any(test_array))
```

```
print (all(test_array))
```

```
OUT:
```

```
False
```

```
False
```

```
True
```

```
True
```

```
[False, False, True]
```

```
True
```

```
False
```

1.11 Loops and iterating

for loops can be used to step through lists, tuples, and other 'iterable' objects.

Iterating through a list:

```
for item in [10,25,50,75,100]:  
    print (item, item**2)
```

OUT:

```
10 100  
25 625  
50 2500  
75 5625  
100 10000
```

A *for* loop may be used to generate and loop through a sequence of numbers (note that a 'range' does not include the maximum value specified):

```
for i in range(100,150,10):  
    print(i)
```

OUT:

```
100  
110  
120  
130  
140
```

A *for* loop may be used to loop through an index of positions in a list:

```
my_list = ['Frodo','Bilbo','Gandalf','Gimli','Sauron']  
  
for i in range(len(my_list)):  
    print ('Index:',i,' Value',my_list[i])
```

OUT:

```
Index: 0 , Value Frodo  
Index: 1 , Value Bilbo  
Index: 2 , Value Gandalf  
Index: 3 , Value Gimli  
Index: 4 , Value Sauron
```

1.11.1 Breaking out of loops or continuing the loop without action

Though it may not be considered best coding practice, it is possible to prematurely escape a loop with the *break* command:

```
for i in range(10): # This loop would normally go from 0 to 9  
    if i == 5:  
        break  
    else:  
        print(i)  
  
print ('Loop complete')
```

OUT:

```
0
1
2
3
4
Loop complete
```

Or, rather than breaking out of a loop, it is possible to effectively skip an iteration of a loop with the *continue* command. This may be placed anywhere in the loop and returns the focus to the start of the loop.

```
for i in range (10):
    if i%2 == 0: # This is the integer remainder after dividing i by 2
        continue
    else:
        print (i)

print ('Loop complete')
```

OUT:

```
1
3
5
7
9
Loop complete
```

1.11.2 Using pass to replace active code in a loop

The *pass* command is most useful as a place holder to allow a loop to be built and have contents added later.

```
for i in range (10):
    # Some code will be added here later
    pass

print ('Loop complete')
```

OUT:

```
Loop complete
```

1.12 List comprehensions: one line loops

List comprehensions are a very Pythonic way of condensing loops and action into a single line.

Here is the long form of a loop:

```
my_list=[1,4,5,7,9,10,3,2,16]
my_new_list=[]

for x in my_list:
    if x>5:
        my_new_list.append(x**2)

print (my_new_list)
```

OUT:

```
[49, 81, 100, 256]
```

This may be done by a single line list comprehension:

```
y = [x**2 for x in my_list if x>5]

print (y)
```

OUT:

```
[49, 81, 100, 256]
```

1.13 Else after while

The *else* clause is only executed when your while condition becomes false. If you *break* out of the loop it won't be executed.

The following loop will run the else code:

```
x = 0
while x < 5:
    x += 1
    print (x)
else:
    print ('Else statement has run')
print ('End')
```

OUT:

```
x = 0

while x < 5:

    x += 1

    print (x)

else:

    print ('Else statement has run')

print ('End')
```

OUT:

```
1
2
3
4
5
Else statement has run
End
```

The following loop with a *break* will not run the *else* code:

```
x = 0
while x < 5:
    x += 1
    print (x)
    if x > 3:
        break
else:
    print ('Else statement has run')
print ('End')
```

OUT:

```
1
2
3
4
End
```

1.14 try ... except (where code might fail)

In perfect code all eventualities will be anticipated and code written for all eventualities. Sometimes though we might want to allow for more general failure of code and provide an alternative route.

In the following code we try to perform some code on a list. For the code to work the list must be long enough and the number must be positive. We use try except to perform the calculation where possible and return a null value if not.

In Python an error is called an 'exception'.

```
import math

x = [1, 2, -4]

# Try to take the square root of elements of a list.
# As a user try entering a list index too high (3 or higher)
# or try to enter index 2 (which has a negative value, -4)

test = int(input('Index position?: '))
try:
    print (math.sqrt(x[test]))
except:
    print ('Sorry, no can do')
```

OUT:

```
Index position?: 2
Sorry, no can do
```

Try except may be used in final code (but it is better to allow for all specific circumstances) or may be used during debugging to evaluate variables at the time of code failure

1.15 Decimal places in output

Python allows a lot of control over formatting of output. But here we will just look at controlling decimal places of output.

There are some different ways. Here is perhaps the most common (because it is most similar to other languages).

The number use is represented in the print function as `%x.yf` where `x` is the total number of spaces to use (if defined this will add padding to short numbers to align numbers on different lines, for example), and `y` is the number of decimal places. `f` informs python that we are formatting a float (a decimal number). The `%x.yf` is used as a placeholder in the output (multiple placeholders are possible) and the values are given at the end of the print statement as below:

```
import math
pi = math.pi
pi_square = pi**2

print('Pi is %.3f, and Pi squared is %.3f' %(pi,pi_square))
```

OUT:

```
Pi is 3.142, and Pi squared is 9.870
```

It is also possible to round numbers before printing (or sending to a file). If taking this approach be aware that this may limit the precision of further work using these numbers:

```
import math
pi = math.pi
pi = round(pi,3)
print (pi)
```

OUT:

```
3.142
```

1.16 Writing to and reading from files

For data analytics we will usually be using two libraries called NumPy and Pandas which have their own simple methods for importing data, but here is the standard Python method. This standard method may sometimes have an advantage that data may be written and read one line at a time without having all data held in memory. This may be useful for very large data files.

1.16.1 File access modes

Files may be opened with different access control models

Text file access modes:

"r" Read (if file does not exist there will be an error)

"w" Write If the file exists contents are overwritten. If the file does not exist, it is created.

"a" Append - will add to file, or will create a file if no file exists

"r+" Read from and write to a file

"w+" Write to and read from a text file. If the file exists contents are overwritten. If the file does not exist, it is created.

"a+" Append and read from a text file. If the file exists, new data is appended to it. If the files doesn't exist, it's created.

1.16.2 Creating a file with write

This method where we open a file to save to with the 'w' argument will create a new file or overwrite an old one.

In this standard way of using a file we begin by opening a file and end by closing the file:

```
text_file = open('write_it.txt', 'w')
text_file.write('Line 1\n')
text_file.write('This is line 2\n')
text_file.write('That makes this line 3\n')
text_file.close()
\begin{verbatim}
```

An alternative to opne and closing the file is using the with statement:

```
\begin{verbatim}
with open('write_it.txt', 'w') as text_file:
    text_file.write('Line 1\n')
    text_file.write('This is line 2\n')
    text_file.write('That makes this line 3\n')
```

1.16.3 Adding to a file

Using the 'a' mode we will add to an existing file. We will use the with statement below, but the open/close method will work as well.

```
with open('write_it.txt', 'a') as text_file:
    text_file.write('This is a fourth line\n')
```


1.16.4 Writing multiple lines

Multiple lines may be written using a loop with the *write* method, or may we can use *writelines*:

```
'write_it.txt'
text_file = open('write_it.txt', 'w')
lines = ['Line 1\n',
        'This is line 2\n',
        'That makes this line 3\n']
text_file.writelines(lines)
text_file.close()
```

1.16.5 Reading a file

We can read a file in one line at a time (useful for processing very large files which cannot be held in memory), or we can read in all lines.

Here we read one line at a time:

```
text_file = open('write_it.txt', 'r')
text_file = open('write_it.txt', 'r')
print(text_file.readline())
print(text_file.readline())
print(text_file.readline())
text_file.close()
```

OUT:

Line 1

This is line 2

That makes this line 3

Or we can read the entire file in as a list:

```
text_file = open('write_it.txt', 'r')
lines = text_file.readlines()
text_file.close()
print (lines)
```

OUT:

```
['Line 1\n', 'This is line 2\n', 'That makes this line 3\n']
\end{verbatim}
```

A text file may be stepped through as part of a loop:

```
\begin{verbatim}
text_file = open('write_it.txt', 'r')
for line in text_file:
    print(line)
text_file.close()
OUT:
```

Line 1

This is line 2

That makes this line 3

1.17 Functions

Functions are (or should be) short sections of code that may be called from within the main programme, or from another function. They usually take some inputs (arguments) and return one or more results.

Functions may be used for two purposes:

- 1) To avoid repeating code when the same code needs to be run in two or more locations in the main programme.
- 2) To structure the code, breaking the code down into smaller pieces. In functional programming the whole programme would be broken down into functions

If a function is more than 10 lines long you may want to think whether it could be broken down further, with some of the code pulled out as a separate function.

Ideally a function should do just one thing; it should be as simple as possible, and easy for another person to follow.

A docstring may be used at the start of a function to provide help to the user - in this case if the user types `help(function_name)` the docstring will be returned. These docstrings may also give examples of the function.

Any variables defined within functions are destroyed when the function closes. Functions are not usually used to change global variables (those variables defined in the main programme code). If a global variable is passed to a function any changes to that variable made by the function are lost when the function ends (and so global variables are protected from being changed by functions). There is a way to force the function to change the global variable, which we will cover in a later lesson, but this should usually be avoided if possible.

1.17.1 Defining a function

A function is defined with the `def` statement. It is followed by the function's inputs (arguments or 'args') which may also be given a default value. The user may enter the inputs either by name, or in the order in which the function expects them.

It is good practice to use verbs to define functions, to describe what they do.

Let's do a very simple example, raising a number to the power of another number, and we will set a default where, if no other number is given, we will raise to the power of zero (which always returns a value of 1). Note that we define a default value with the `=` sign.

```
def raise_to_power (number, power = 0):  
    result = number ** power  
    return result
```

Once the function has been defined we can call it:

```
print (raise_to_power(2,6))
```

OUT:

64

We could describe the inputs by their names if we wanted to:

```
print (raise_to_power(number = 2, power = 6))
```

OUT:

64

If we define the inputs by name, the order becomes unimportant (without the names the function expects inputs in the order defined in the function):

```
print (raise_to_power(power = 6, number = 2))
```

OUT:

64

And remember that we set a default of power to zero, so if we only pass the first argument ('number'), the power defaults to zero:

```
print (raise_to_power(2))
```

OUT:

1

1.17.2 Adding help to a function with a docstring

A help docstring may be added to a function using triple quotes (' or ") after the function has been defined. This may be called later by the user using help.

```
def raise_to_power (number, power = 0):
    """
    This functions takes two numbers, and raises the first number to the power of the second.
    If no second number, or power, is given then power defaults to zero
    """

    result = number ** power
    return result
```

Calling help:

```
help (raise_to_power)
```

OUT:

Help on function raise_to_power in module __main__:

```
raise_to_power(number, power=0)
    This functions takes two numbers, and raises the first number to the power of the second.
    If no second number, or power, is given then power defaults to zero
```

1.17.3 Returning two or more results

Sometimes we may wish to return more than one value. There are two ways of doing this:

1) Return the values separately (the code that calls the argument must also refer to both results:

```
def calculate_sum_and_product(a, b):
    """Return sum and product of two numbers"""

    my_sum = a + b
    my_product = a * b
    return my_sum, my_product

# When calling the function we need to put both results generated by the function

calculated_sum, calculated_product = calculate_sum_and_product(5,6)
print (calculated_sum)
print (calculated_product)
```

OUT:

11
30

2) A more common way would be to return a single result that is a container (commonly a list or a tuple) of the two results. Here we return a tuple of the two results (remember that tuples are like lists, but they cannot be changed). We can tell that the returned value is a tuple by the use of curved brackets (a list could also be used, which we would spot by square brackets).

```
def calculate_sum_and_product(a, b):  
    """Return sum and product of two numbers"""  
    my_sum = a + b  
    my_product = a * b  
    result = (my_sum, my_product) # this is a tuple  
    return results  
  
results = calculate_sum_and_product(5,6)  
  
# Individual results are obtained by their index:  
  
print (results[0])  
print (results[1])
```

OUT:

11
30

1.18 Lambda functions (one line functions), and map/filter/reduce

Lambda functions are a Pythonic way of writing simple functions in a single line.

They may be applied to a list using the *map* statement.

The function *filter* offers a concise way to filter out all the elements of a list based on the results of a function.

The function *reduce* continually applies a function to a sequence until a single value is returned.

1.18.1 Defining one-line lambda functions

Instead of defining the function with `def`, we define the function with *lambda*.

Here is a 'normal' function:

```
# A named function that returns the power of a number:
```

```
def f(x, y):  
    result = x ** y  
    return result
```

```
print (f(4,3))
```

OUT:

64

Re-writing as a lambda function:

```
g = lambda x,y: x**y
```

```
print (g(4,3))
```

OUT:

64

1.18.2 Applying a lambda function to a list with map

In the example below we define a lambda function to convert Celsius to Fahrenheit, and then apply it to a list.

We will first define lambda separately, and then map, before looking at how they can be combined a single line.

The *map* command generates a 'map object'. To return a list we need to convert that to a list as below

```
celsius = [0, 10, 20, 30]  
g = lambda x: (9/5)*x + 32  
fahrenheit = list(map(g, celsius))
```

```
print (fahrenheit)
```

OUT:

[32.0, 50.0, 68.0, 86.0]

This may be compressed further into one line. But there is always a balance to be had between being as concise as possible, and being as clear for others as possible. It may, however, generally be assumed that other experienced Python coders will be familiar with a combined one-line lambda/map function.

```
celsius = [0, 10, 20, 30]
fahrenheit = list(map(lambda x: (9/5)*x + 32, celsius))

print (fahrenheit)

[32.0, 50.0, 68.0, 86.0]
```

1.18.3 Filtering a list with a lambda function

The function *filter* offers a concise way to filter out all the elements of a list, for which the function returns True. The function `filter(f,l)` needs a function `f` as its first argument. `f` returns a Boolean value, i.e. either True or False.

This function will be applied to every element of the list `l`. Only if `f` returns True will the element of the list be included in the result list.

Here is an example where `filter` is used to return even numbers from a list. As with the `map` example above we will combine the `filter` and `lambda` function in a single line.

```
my_list = [0,1,2,3,4,5,6,7,8,9,10]
result = list(filter(lambda x: x % 2 == 0, my_list))

print (result)

OUT:

[0, 2, 4, 6, 8, 10]
```

1.18.4 Reducing a list with a repeated lambda function call

The `reduce` function will continually apply a lambda function until only a single value is returned. Though it used to be in core python, it is now found in the `functools` module.

A simple example is to use `reduce` to multiply all values in a list. The `reduce` function replaces the first two elements of a list with the product of those two elements. It will continue until one value is left.

```
import functools
my_list = [10,30,34,56,89]
result = functools.reduce(lambda x,y: x*y, my_list)

print (result)

OUT:

50836800
```

1.18.5 Alternatives to map, filter and reduce

Most results from `map/filter` and `reduce` may also be obtained from using list comprehensions, as previously outlined. Some see list comprehensions as more Pythonic than `map`, `filter` and `reduce`.

1.19 Accessing date and time, and timing code

1.19.1 Accessing date and time

It may sometimes be useful to access the current date and/or time. As an example, when writing code it may be useful to access the time at particular stages to monitor how long different parts of code are taking.

To access date and time we will use the *datetime* module (which is held in a package that is also, a little confusingly called *datetime*):

```
from datetime import datetime
```

To access the current date and time we can use the `now` method:

```
print (datetime.now())
```

OUT:

```
2018-03-28 08:54:13.623198
```

You might not have expected to get the time to the nearest microsecond! But that may be useful at times when timing short bits of code (which may have to run many times).

But we can access the date and time in different ways:

```
current_datetime = datetime.now()

print ('Date:', current_datetime.date()) # note the () after date
print ('Year:', current_datetime.year)
print ('Month:', current_datetime.month)
print ('Day:', current_datetime.day)
print ('Time:', current_datetime.time()) # note the () after time
print ('Hour:', current_datetime.hour)
print ('Minute:', current_datetime.minute)
print ('Second:', current_datetime.second)
```

OUT:

```
Date: 2018-03-28
Year: 2018
Month: 3
Day: 28
Time: 08:54:13.636654
Hour: 8
Minute: 54
Second: 13
```

Having microseconds in the time may look a little clumsy. So let's format the date and using the *strftime* method.

```
now = datetime.now()
print (now.strftime('%Y-%m-%d'))
print (now.strftime('%H:%M:%S'))
```

OUT:

```
2018-03-28
08:54:13
```

1.19.2 Timing code

We can use `datetime` to record the time before and after some code to calculate the time taken, but it is a little simpler to use the `time` module which keeps all time in seconds (from January 1st 1970) and gives even more accuracy to the time:

```
import time

time_start = time.time()

for i in range(100000):
    x = i ** 2

time_end = time.time()
time_taken = time_end - time_start

print ('Time taken:', time_taken)
```

OUT:

Time taken: 0.03149104118347168

Chapter 2

NumPy and Pandas

2.1 Numpy and Pandas

For much of healthcare data analytics and modelling we will use NumPy and Pandas as the key containers of our data. These libraries allow efficient manipulation and analysis of very large data sets. They are very considerably faster than using a 'pure Python' approach.

NumPy and Pandas are distributed with all main scientific Python distributions.

2.1.1 NumPy

NumPy is a library for supporting work with large multi-dimensional arrays, with many mathematical functions. NumPy is compatible with many other libraries such as Pandas (see below), Matplotlib (for plotting), and many Python maths, stats and optimisation libraries.

When we import NumPy as a library it is standard practice to use the *as* statement which allows it to be referenced with a shorter name. We will import as *np*:

```
import numpy as np
```

2.1.2 Pandas

Pandas is a library that allows manipulation of large arrays of data. Data may be indexed and manipulated based on index. Data is readily pivoted, reshaped, grouped, merged.

We will use pandas alongside numpy. Generally, NumPy is faster for mathematical functions, but Pandas is more powerful for data manipulation.

As with numpy, we import with a shortened name:

```
import pandas as pd
```

2.2 Numpy basics: building an array from lists, basic statistics, converting to booleans, referencing the array, and taking slices

Most commonly we will be loading files into NumPy arrays, but here we build an array from lists and perform some basic stats on the array.

The examples below construct and use a 2 dimensional array (which may be thought of as 'rows and columns'). Later we will look at higher dimensional arrays.

2.2.1 Building an array from lists

We use the `np.array` function to build an array from existing lists. Here each list represents a row of a data table.

```
import numpy as np

row0 = [23, 89, 100]
row1 = [10, 51, 99]
row2 = [40, 78, 102]
row3 = [35, 81, 110]
row4 = [50, 75, 95]
row5 = [65, 51, 101]

data = np.array([row0, row1, row2, row3, row4, row5])
```

We now have a data array:

```
print (data)
```

OUT:

```
[[ 23  89 100]
 [ 10  51  99]
 [ 40  78 102]
 [ 35  81 110]
 [ 50  75  95]
 [ 65  51 101]]
```

2.2.2 Performing basic statistics on the array

We can see, for example, the mean of all the data

```
print (data.mean())
```

OUT:

```
69.72222222222223
```

Or we can use the *axis* argument to show the mean by column or mean by row:

```
print (np.mean(data,axis=0)) # average all rows in each column
print (np.mean(data,axis=1)) # average all columns in each row
```

OUT:

```
[ 37.16666667  70.83333333 101.16666667]
[70.66666667  53.33333333 73.33333333 75.33333333 73.33333333 72.33333333]
```

Some other commonly used statistics (all of which are by column, or dimension 0, below):

```
print ('Mean:\n', np.mean(data, axis=0))
print ('Median:\n', np.median(data, axis=0))
print ('Sum:\n', np.sum(data, axis=0))
print ('Maximum:\n', np.max(data,axis=0))
print ('Minimum\n:', np.min(data,axis= 0))
print ('Range:\n', np.ptp(data, axis=0))
print ('10th percentile:\n', np.percentile(data, 10, axis=0))
print ('90th percentile:\n', np.percentile(data, 90, axis=0))
print ('Population standard deviation\n', np.std(data, axis=0))
print ('Sample standard deviation:\n', np.std(data, axis=0, ddof=1))
print ('Population variance:\n', np.var(data, axis=0))
print ('Sample variance:\n', np.var(data, axis=0, ddof=1))
```

OUT:

```
Mean:
 [ 37.16666667  70.83333333 101.16666667]
Median:
 [ 37.5  76.5 100.5]
Sum:
 [223 425 607]
Maximum:
 [ 65  89 110]
Minimum
: [10 51 95]
Range:
 [55 38 15]
10th percentile:
 [16.5 51.  97. ]
90th percentile:
 [ 57.5  85. 106. ]
Population standard deviation
 [17.75215167 14.6562463  4.52462399]
Sample standard deviation:
 [19.44650783 16.05511341  4.95647724]
Population variance:
 [315.13888889 214.80555556  20.47222222]
Sample variance:
 [378.16666667 257.76666667  24.56666667]
```

The returned array may be referenced by index number (beginning at zero):

```
results = np.mean(data, axis=0)
```

```
print (results[0])
```

OUT:

```
37.166666666666664
```

Basic python may also be incorporated into the statistics. For example the 10th percentiles, from 0 to 100 may be calculated in a loop (remember that the standard Python range function goes up to, but does not include, the maximum value given, so we need to put a higher maximum than 100 in order to include the 100th percentile in the loop):

```
for percent in range(0,101,10):
    print(percent,'percentile:',np.percentile(data, percent, axis=0))
```

OUT:

```

0 percentile: [10. 51. 95.]
10 percentile: [16.5 51. 97. ]
20 percentile: [23. 51. 99.]
30 percentile: [29. 63. 99.5]
40 percentile: [ 35. 75. 100.]
50 percentile: [ 37.5 76.5 100.5]
60 percentile: [ 40. 78. 101.]
70 percentile: [ 45. 79.5 101.5]
80 percentile: [ 50. 81. 102.]
90 percentile: [ 57.5 85. 106. ]
100 percentile: [ 65. 89. 110.]

```

2.2.3 Converting values to a boolean True/False

We can use test values against some standard. For example, here we test whether a value is equal to, or greater than, the mean of the column.

```

column_mean = np.mean(data, axis=0)
greater_than_mean = data >= column_mean

```

```
print (greater_than_mean)
```

OUT:

```

[[False  True False]
 [False False False]
 [ True  True  True]
 [False  True  True]
 [ True  True False]
 [ True False False]]

```

True/False values in Python may also be used in calculations where False has an equivalent value to zero, and True has an equivalent value to 1.

```
print (np.mean(greater_than_mean, axis=0))
```

OUT:

```
[0.5          0.66666667 0.33333333]
```

2.2.4 Referencing arrays and taking slices

NumPy arrays are referenced similar to lists, except we have two (or more dimensions). For a two dimensional array the reference is [dimension_0, dimension_1], which is equivalent to [row, column].

REMEMBER: Like lists, indices start at index zero, and when taking a slice the slice goes up to, but does not include, the maximum value given.

```

print ('Row zero:\n', data[0,:])
print ('Column one:\n', data[:,1])
# Note in the example below a slice goes up to ,but dot include, the maximum index
print ('Rows 0-3, Column zero:\n', data[0:4,0])
print ('Row 1, Column 2\n', data[1,2])

```

OUT:

```

Row zero:
 [ 23  89 100]
Column one:

```

```
[89 51 78 81 75 51]
Rows 0-3, Column zero:
[23 10 40 35]
Row 1, Column 2
99
```